Contents

Introduction

Advances in quantum computing will significantly increase the range of problems that can be attacked with computation. The most important point for cybersecurity is not that quantum computers will make every problem easy. They will not. The point is that a small set of quantum algorithms changes the cost of solving specific mathematical problems that are used to secure almost every modern network. Public-key cryptography depends on mathematical asymmetry: it is easy to multiply two large primes, but classically difficult to recover those primes from their product; it is easy to exponentiate in a finite group, but classically difficult to solve the corresponding discrete logarithm. A sufficiently capable quantum computer can alter that balance.

This expanded chapter keeps the cybersecurity emphasis of the original discussion while adding more detail about what a qubit is, how a physical qubit is built, and why different companies and national laboratories are pursuing different machine architectures. It also separates several ideas that are often confused in popular accounts: a qubit is not simply a bit that is both 0 and 1; a high qubit count is not automatically better; quantum advantage demonstrations are not the same thing as general-purpose, fault-tolerant computing; and a quantum annealer is not identical to a gate-model quantum computer.

The practical consequence is that organizations should plan before large cryptographically relevant quantum computers exist. Data can be collected now and decrypted later, a strategy often called harvest now, decrypt later. Long-lived secrets, code-signing keys, identity credentials, medical records, government archives, financial records, and industrial intellectual property are therefore exposed to a future capability even if today's attackers do not yet have that capability. Quantum readiness is thus a risk-management problem, not only a physics problem.

Quantum computers are still engineering-limited. Current systems are noisy, specialized, and expensive to operate. They require careful shielding, precise control electronics, specialized fabrication, and complex calibration. Even so, the field has moved rapidly. IBM, Google, university groups in China, IonQ, Quantinuum, D-Wave, Xanadu, QuEra, Atom Computing, Microsoft, and other organizations have announced systems or roadmaps that show clear progress in scale, fidelity, error correction, cloud access, and specialized demonstrations. The central question is no longer whether quantum information can be processed; it is how soon large-scale, reliable, useful quantum computation can be engineered.

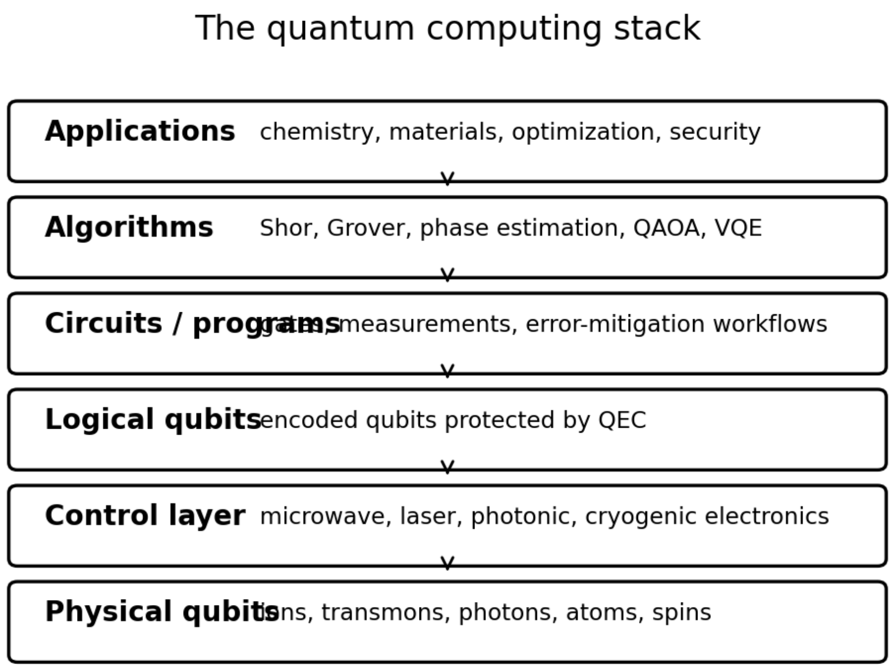

Figure 1 Quantum computing is best understood as a stack that begins with physical qubits and ends with applications.

What this means for cryptography

Public-key cryptography is the most visible cybersecurity issue because widely deployed algorithms such as RSA, finite-field Diffie-Hellman, elliptic-curve Diffie-Hellman, and elliptic-curve digital signatures depend on problems for which efficient quantum algorithms are known. Peter Shor developed a quantum algorithm that factors integers and solves discrete logarithms in polynomial time in the input size. In ordinary terms, this means that once a sufficiently large, error-corrected quantum computer exists, the mathematical problems that protect RSA and Diffie-Hellman style systems will no longer provide the intended security margin.

The key phrase is sufficiently large and error-corrected. A real attack on modern RSA or elliptic-curve keys is not expected from a small noisy processor. The problem requires logical qubits that can survive long computations, a large number of reliable gates, and enough error correction to keep the computation coherent. That is why quantum risk planning focuses on migration timelines. It is possible for a technology to be years away from breaking a key directly while still being close enough to justify replacing protocols in long-lived systems.

Symmetric cryptography is affected differently. Grover's algorithm gives a quadratic speedup for unstructured search, so it reduces the brute-force security of an n-bit key to roughly n/2 bits against an ideal quantum adversary. That is serious but manageable. AES-128 should not be treated as providing a 128-bit margin against a large quantum adversary, but AES-256 remains a practical countermeasure for most applications. Hash functions and message authentication codes also need attention, but the most urgent migration problem is public-key establishment and digital signatures.

The National Institute of Standards and Technology finalized its first three post-quantum cryptography standards in 2024: FIPS 203 for ML-KEM, derived from CRYSTALS-Kyber, for key establishment; FIPS 204 for ML-DSA, derived from CRYSTALS-Dilithium, for digital signatures; and FIPS 205 for SLH-DSA, derived from SPHINCS+, as a stateless hash-based signature option. NIST has also indicated that a Falcon-derived signature standard, FN-DSA, is planned. These standards do not require organizations to wait for quantum computers. They are intended to be deployed before the threat becomes operational (NIST, 2024).

The cryptographic migration will not be a simple drop-in replacement. Protocols, hardware security modules, certificates, embedded devices, firmware update systems, secure boot chains, VPNs, TLS deployments, PKI governance, and long-term archives all need evaluation. Hybrid key exchange, in which classical and post-quantum mechanisms are combined, is one bridge strategy because it can protect against future quantum attacks without immediately abandoning mature classical mechanisms. Over time, however, organizations will need native post-quantum cryptographic agility: the ability to change algorithms, key sizes, and parameters without redesigning entire systems.

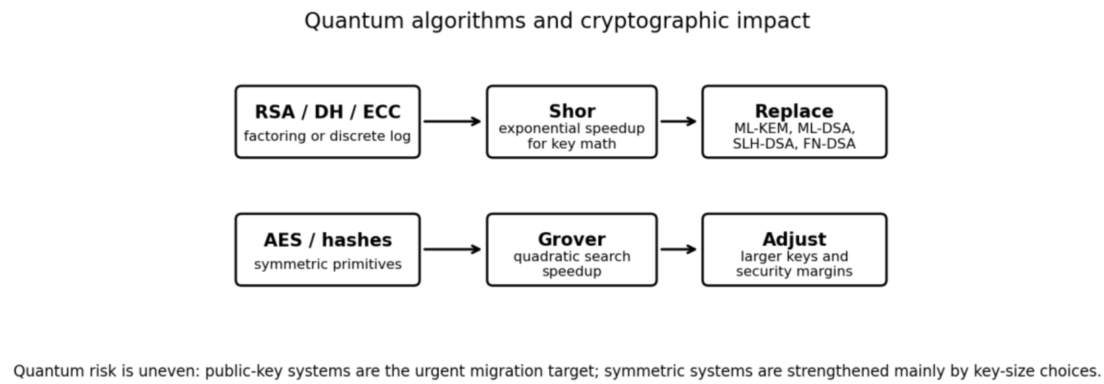

Figure 2 Quantum algorithms affect asymmetric and symmetric cryptography in different ways.

Table 1 Quantum impact on major cryptographic primitives.

Primitive	Underlying security idea	Quantum relevance	Migration response
RSA encryption and signatures	Integer factorization	Shor algorithm	Replace with approved post-quantum key establishment and signatures.
Finite-field Diffie-Hellman	Discrete logarithm	Shor algorithm	Replace key exchange with ML-KEM or a hybrid scheme during transition.
Elliptic-curve Diffie-Hellman and ECDSA	Elliptic-curve discrete logarithm	Shor algorithm	Replace key agreement and signatures; review certificates and identity systems.
AES and other symmetric ciphers	Exhaustive key search	Grover algorithm	Use longer keys where appropriate; AES-256 is the normal conservative choice.
Hash functions	Preimage and collision search	Grover-like and collision algorithms	Use sufficiently large output sizes and modern constructions.

A brief history of quantum computing

A useful history of quantum computing begins with the realization that nature itself is quantum. Richard Feynman argued in the early 1980s that simulating quantum physics on classical computers was inherently inefficient and that a computer governed by quantum mechanics might simulate quantum systems more naturally. David Deutsch later formalized universal quantum computation. In 1994 Shor showed that a quantum computer could factor large integers and compute discrete logarithms efficiently, making quantum computing immediately relevant to cryptography. In 1996 Grover showed a quantum search speedup that also has cryptographic implications.

Early laboratory demonstrations used nuclear magnetic resonance, photons, trapped ions, and small superconducting circuits. These experiments were extremely small by today's standards, but they proved that quantum information could be initialized, manipulated, entangled, and measured. Through the 2000s and 2010s the field moved from proof-of-principle devices to cloud-accessible processors, commercial roadmaps, and specialized demonstrations of quantum advantage. The most visible milestone was Google's 2019 random circuit sampling experiment with Sycamore, which claimed a dramatic speedup for a narrowly defined sampling task. Chinese groups soon reported their own superconducting and photonic sampling demonstrations, and the debate turned from whether quantum devices can outperform classical machines on some tasks to how useful and robust those tasks are.

The 2020s have been marked by competition among hardware modalities. Superconducting processors led by IBM, Google, and Chinese university groups have scaled to hundreds or more physical qubits. Trapped-ion systems led by Quantinuum and IonQ have emphasized fidelity, connectivity, and mid-circuit measurement. Photonic systems from USTC and Xanadu have shown large boson-sampling experiments. Neutral-atom systems from QuEra and Atom Computing have highlighted rapid physical scaling and reconfigurable arrays. D-Wave has continued to develop quantum annealing systems with thousands of superconducting qubits for optimization-style problems. These paths are not interchangeable, but each contributes important engineering ideas.

The next historical phase is expected to be fault tolerance. Noisy intermediate-scale quantum processors can be scientifically valuable, but they are not enough for large cryptanalytic attacks or the most demanding chemistry simulations. The important milestones now involve logical qubits, surface-code or alternative quantum error correction, better calibration, lower error rates, higher connectivity, modularity, and sustained computations measured in millions or billions of reliable operations rather than isolated demonstrations.

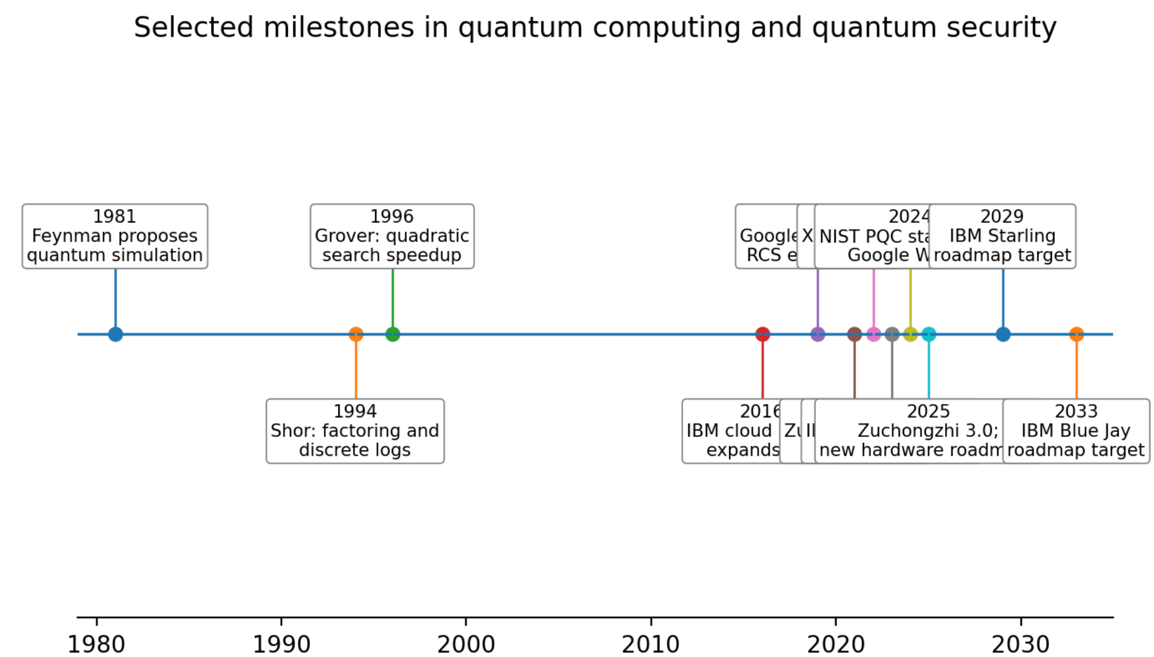

Figure 3 Selected milestones in quantum computing and quantum security. Roadmap targets are not guarantees.

Table 2 Selected milestones in context.

Period	Milestone	Why it matters
1981-1982	Feynman argues that quantum systems should be simulated by quantum devices.	Motivates quantum simulation as a native application.
1994	Shor publishes factoring and discrete-log quantum algorithms.	Creates the central public-key cryptography concern.
1996	Grover publishes a quantum search algorithm.	Shows a quadratic speedup relevant to symmetric key sizes.
2016	Cloud access to quantum processors expands.	Allows broader experimentation by researchers and students.
2019	Google reports the Sycamore random circuit sampling experiment.	Popularizes the term quantum supremacy or quantum advantage.
2021-2022	USTC reports Zuchongzhi and Jiuzhang 2.0; Xanadu reports Borealis.	Demonstrates multiple paths to special-purpose advantage claims.
2023-2024	IBM announces Condor, Heron, System Two; Google announces Willow; NIST finalizes first PQC standards.	Shows simultaneous progress in hardware and defensive cryptography.
2029-2033	IBM roadmap targets Starling and Blue Jay fault-tolerant systems.	Highlights the industry shift from physical scale toward logical operations.

What is a Quantum Computer?

A quantum computer is a machine that stores and processes information using controllable quantum systems. The system might be a superconducting circuit, an ion held in an electromagnetic trap, a photon traveling through an optical circuit, a neutral atom held by optical tweezers, an electron spin in a semiconductor quantum dot, a defect in diamond, or another physical platform. What makes the system quantum is not its size or its vendor; it is the use of quantum states, reversible quantum gates or controlled evolution, measurement, and interference to process information.

A classical digital computer stores information as bits. A bit has a value of 0 or 1, and a classical logic gate maps one set of bit values to another. A quantum computer stores information in qubits. A qubit can be in a superposition of computational basis states, and a multi-qubit register can hold amplitudes over many possible basis states. Quantum gates change those amplitudes. A measurement returns ordinary classical information, but the probability of each outcome depends on how the amplitudes have been prepared, interfered, and entangled.

This last point is essential. Quantum computing does not provide a shortcut by simply trying every answer at once and reading them all out. When a quantum register is measured, it yields a single classical outcome or a distribution of outcomes after many runs. Quantum algorithms are powerful only when they exploit structure. They arrange amplitudes so that wrong answers tend to interfere destructively and useful answers tend to interfere constructively. Shor's algorithm uses the structure of periodicity; Grover's algorithm amplifies marked states; quantum simulation uses the fact that quantum systems naturally evolve in a quantum state space.

Physical quantum computers are therefore hybrid machines. A cryostat, vacuum chamber, laser system, microwave control chain, photonic chip, or other apparatus hosts the qubits. Classical control computers compile programs into pulses or optical settings. Measurement electronics convert fragile quantum information into classical bits. Error mitigation, calibration, and error correction software then interpret or stabilize the result. The machine is quantum at the information-processing core, but it is surrounded by classical engineering.

The dominant model for universal quantum computation is the gate model. In the gate model, a program is a circuit made of one-qubit and two-qubit operations followed by measurements. Another important model is measurement-based quantum computing, where an entangled resource state is prepared and computation is driven by measurements. Quantum annealing is different: it encodes optimization problems into an energy landscape and tries to reach a low-energy configuration. Annealing can be useful for certain optimization formulations, but it should not be treated as equivalent to a gate-model processor capable of running Shor's algorithm as normally described.

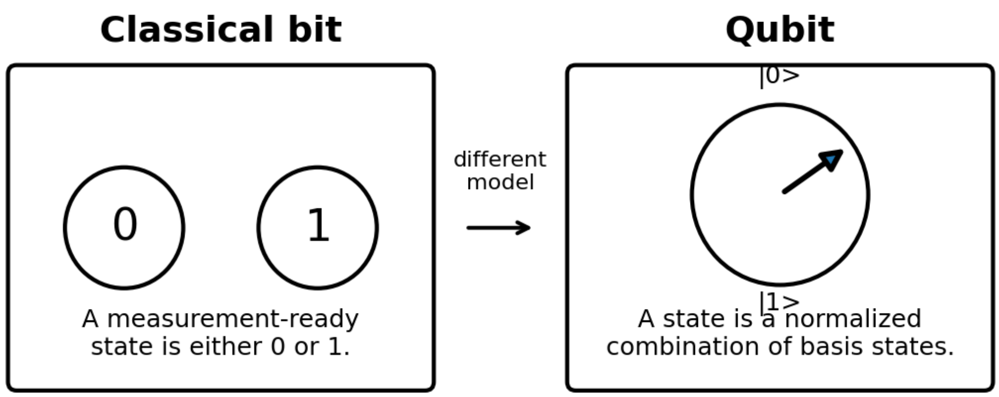

Figure 4 Classical bits and qubits represent information with different mathematical models.

What is a Qubit?

A qubit is the basic unit of quantum information. Mathematically, it is a normalized vector in a two-dimensional complex vector space. The two standard computational basis states are written |0> and |1>. A general pure qubit state can be written as |psi> = alpha|0> + beta|1>, where alpha and beta are complex probability amplitudes and |alpha|^2 + |beta|^2 = 1. If the qubit is measured in the computational basis, the outcome is 0 with probability |alpha|^2 and 1 with probability |beta|^2.

This notation is called Dirac or bra-ket notation. A ket such as |0> represents a vector. In a simple column-vector representation, |0> corresponds to the column vector [1, 0]^T and |1> corresponds to [0, 1]^T. The notation becomes especially useful with many qubits. Two qubits have four basis states: |00>, |01>, |10>, and |11>. Three qubits have eight basis states. In general, n qubits have 2^n computational basis states. A quantum state assigns an amplitude to each basis state, subject to normalization.

The exponential number of amplitudes is the reason quantum systems are hard to simulate classically, but it is not the same as having exponential classical memory that can be read out freely. A measurement of n qubits produces n classical bits. The challenge of algorithm design is to use gates so that the measurement distribution reveals the property of interest. The qubits are manipulated before measurement, and the structure of the algorithm determines whether the useful information is likely to appear.

The Bloch sphere is a common visualization for a single pure qubit. The north pole is |0>, the south pole is |1>, and points on the sphere represent possible pure states. One standard parameterization is |psi> = cos(theta/2)|0> + exp(i phi) sin(theta/2)|1>. The angles theta and phi describe the orientation of the state vector. A global phase does not change measurement probabilities and is not physically observable by itself, but relative phase between components is crucial because it affects interference.

A qubit can be implemented in many physical ways. In a trapped ion, |0> and |1> might be two hyperfine or electronic energy levels of an ion. In a superconducting transmon, they are the two lowest energy states of a nonlinear electrical oscillator. In a photon, they might be two polarizations or two paths. In a neutral atom, they might be two long-lived internal states. The abstract mathematics is the same, but the engineering required to initialize, control, couple, and measure the qubit is very different.

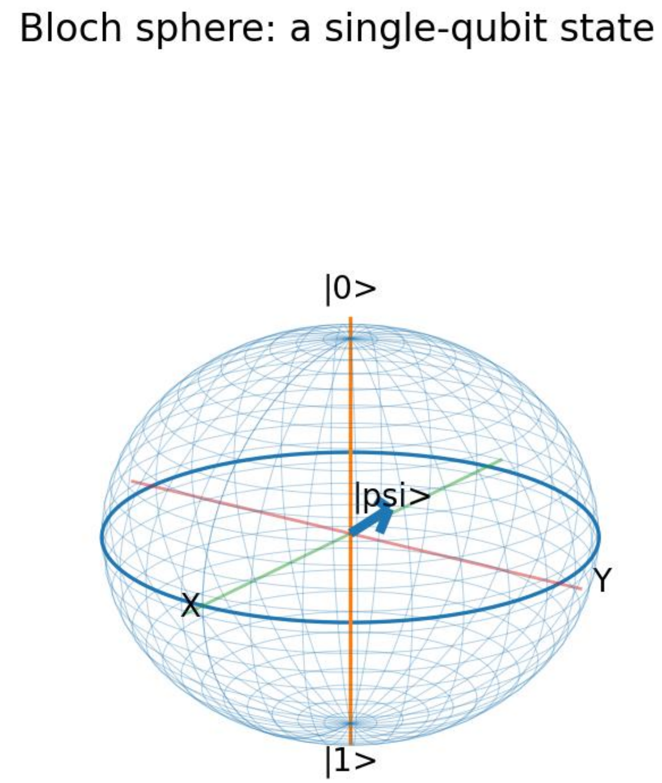

Figure 5 Bloch sphere representation of a single-qubit state. The figure is conceptual; real hardware uses many different physical encodings.

Table 3 Essential qubit vocabulary.

Concept	Notation or idea	Meaning
Computational basis	|0> and |1>	The two reference states used for measurement and logic.
Superposition	|psi> = alpha|0> + beta|1>	The qubit state is a normalized combination of basis states.
Probability amplitudes	alpha and beta	Complex values; squared magnitudes give measurement probabilities.
Normalization	|alpha|^2 + |beta|^2 = 1	The total probability of all measurement outcomes must be one.
Relative phase	phase difference between amplitudes	Controls interference and is essential to quantum algorithms.
Measurement	returns 0 or 1 in a chosen basis	Produces classical data and generally changes the quantum state.
Entanglement	multi-qubit state not separable into single-qubit states	Creates correlations stronger than any classical shared randomness model.

It is also important to distinguish pure states from mixed states. A pure state is described by a state vector. A mixed state describes uncertainty or noise in which the system is not represented by one state vector alone but by a density matrix. Real qubits are often mixed to some degree because of imperfect preparation, environmental coupling, control noise, and measurement error. This is one reason experimental papers report fidelities, coherence times, error rates, and benchmark values rather than just qubit counts.

Another important distinction is between physical qubits and logical qubits. A physical qubit is an actual device-level quantum system. A logical qubit is an encoded qubit protected by quantum error correction. Large-scale algorithms will require logical qubits because physical qubits fail too often over long computations. The conversion from physical qubits to logical qubits is expensive: depending on error rates, code design, and application, one logical qubit may require many physical qubits and repeated syndrome measurements. This overhead is why a device with many physical qubits may still not be capable of the cryptographic attacks that drive post-quantum planning.

Quantum Physics Basics

The quantum world is probabilistic, but not random in the same way as an ordinary coin flip. Quantum mechanics uses amplitudes, and amplitudes can interfere. A particle such as a photon or electron is described by a wave function that encodes probabilities for different measurement outcomes. Before measurement, the state can be a superposition of possibilities. After measurement in a specified basis, a definite outcome is observed and the state is updated accordingly. This is often called wave-function collapse, although interpretations of what collapse means differ.

Wave-particle duality is one of the historical roots of quantum theory. The double-slit experiment shows that particles such as photons and electrons can produce interference patterns associated with waves, even when sent one at a time. If which-path information is obtained, the interference disappears. This behavior is not merely a curiosity; it is the reason phase and interference can carry computational value. Quantum algorithms engineer interference patterns in an abstract state space rather than on a screen behind two slits.

Entanglement is the quantum correlation that has no direct classical counterpart. Two qubits are entangled when the state of the pair cannot be described as independent states for each qubit. For example, a Bell state can be written as (|00> + |11>)/sqrt(2). Measuring one qubit gives a random result, but the other qubit is perfectly correlated when measured in the same basis. Entanglement does not allow faster-than-light communication, but it is a resource for quantum algorithms, quantum teleportation, error correction, and quantum networking.

No-cloning is another important principle. An unknown quantum state cannot be copied perfectly. This creates challenges for error correction because classical computers often protect information by duplicating it. Quantum error correction cannot simply make copies of a qubit. Instead, it encodes information across entangled states of many physical qubits and measures error syndromes that reveal what kind of error occurred without directly revealing the encoded quantum information.

Decoherence is the process by which a quantum system loses its useful phase relationships because it becomes entangled with its environment. In practical terms, the qubit stops behaving as the clean, controlled two-level system intended by the designer. Temperature, electromagnetic radiation, material defects, cosmic rays, crosstalk, laser noise, charge noise, magnetic noise, and imperfect control pulses can all contribute. Decoherence is not only an abstract physics problem; it is an engineering problem that influences every wire, mirror, vacuum seal, chip layout, and calibration routine in a quantum computer.

Different hardware platforms fight decoherence in different ways. Superconducting processors are cooled to millikelvin temperatures in dilution refrigerators. Trapped ions are held in ultrahigh vacuum and manipulated with lasers or microwaves. Photonic processors use particles that interact weakly with the environment but are difficult to make interact with each other. Neutral atoms use laser cooling and optical tweezers. Spin qubits exploit small semiconductor structures and often operate at cryogenic temperatures. The diversity of approaches reflects the fact that no platform has yet solved all scaling, fidelity, connectivity, and manufacturing problems simultaneously.

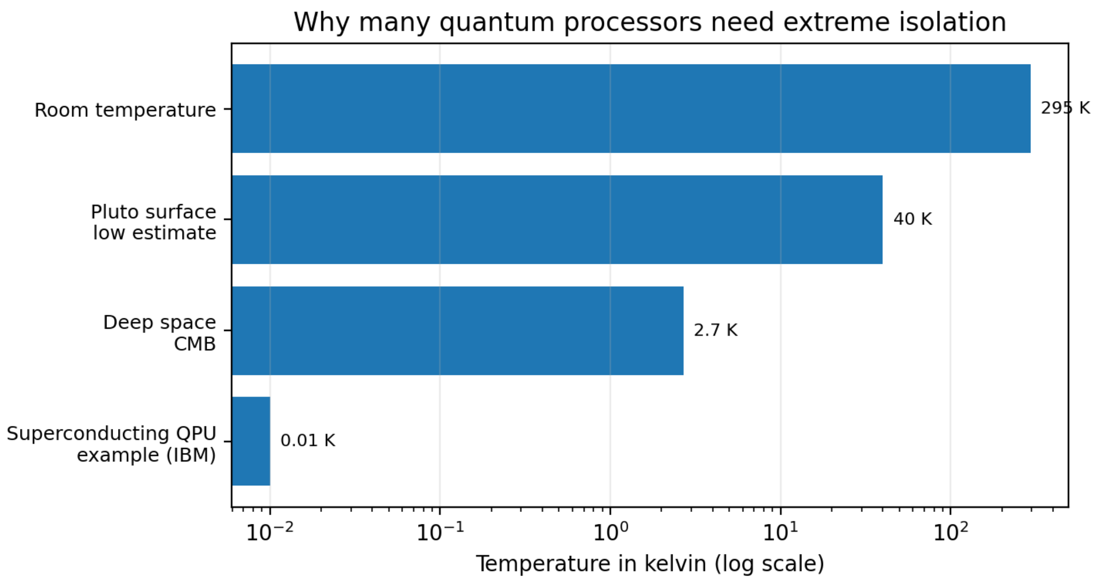

Figure 6 Superconducting quantum processors operate near absolute zero. Source note: IBM describes its processors as operating at about a hundredth of a degree above absolute zero (IBM, 2026).

Quantum gates and circuits

In the gate model, computation is represented as a sequence of quantum gates. A one-qubit gate rotates or changes the state of a single qubit. The X gate is analogous to a bit flip because it maps |0> to |1> and |1> to |0>. The Z gate changes phase. The H, or Hadamard, gate creates equal superpositions and is one of the most common ways to start interference-based circuits. Two-qubit gates such as CNOT, controlled-Z, and iSWAP create entanglement or conditional dynamics between qubits. A universal gate set can approximate any quantum computation to arbitrary accuracy.

Quantum gates are ideally unitary operations, which means they are reversible and preserve total probability. Measurement is different. Measurement maps a quantum state to classical information and is not reversible in the same way. Modern processors increasingly support mid-circuit measurement, reset, and conditional logic. That means a program can measure one qubit during a circuit, use the result to decide later operations, and reuse the qubit. These capabilities are crucial for error correction and more advanced algorithms.

Table 4 Common quantum circuit operations.

<table>
<tr><th>Gate or operation</th><th>Type</th><th>Role in circuits</th></tr>
<tr><td>X</td><td>One-qubit bit flip</td><td>Maps |0> to |1> and |1> to |0>.</td></tr>
<tr><td>Z</td><td>One-qubit phase flip</td><td>Changes the sign of the |1> component.</td></tr>
<tr><td>H</td><td>Hadamard</td><td>Creates equal superpositions and changes measurement basis.</td></tr>
<tr><td>S and T</td><td>Phase gates</td><td>Provide controlled phase rotations; T is important for universal fault-tolerant computation.</td></tr>
<tr><td>CNOT</td><td>Two-qubit controlled operation</td><td>Flips a target qubit conditioned on the control qubit.</td></tr>
<tr><td>CZ</td><td>Two-qubit controlled phase</td><td>Applies a conditional phase and is common in neutral atom and superconducting designs.</td></tr>
<tr><td>Measurement</td><td>Quantum-to-classical readout</td><td>Returns classical bits and is central to error correction.</td></tr>
</table>

Decoherence, noise, and quantum error correction

The main obstacle to useful quantum computing is error. A physical qubit is fragile. It can relax from |1> to |0>, lose phase information, suffer leakage into non-computational states, or be measured incorrectly. Two-qubit gates are usually noisier than one-qubit gates because they require coupling two systems while avoiding unwanted interactions with the rest of the device. When circuits become deep, errors accumulate. This is why simple demonstrations can succeed while long cryptographic algorithms remain out of reach.

Coherence times are usually summarized by T1 and T2. T1 is the relaxation time, often associated with energy decay. T2 is the dephasing time, associated with loss of phase coherence. A qubit with a long T1 and T2 is not automatically a useful qubit, because gate speed, readout fidelity, connectivity, crosstalk, and calibration stability also matter. Conversely, a qubit with shorter coherence can still be useful if gates are very fast and errors are low. The relevant metric is the error per operation and the ability to correct those errors before they accumulate.

Quantum error correction encodes one logical qubit into a larger entangled state of many physical qubits. The system repeatedly measures auxiliary information called syndromes. A syndrome indicates whether an error pattern is likely to have occurred, but it does not reveal the protected quantum information itself. A classical decoder interprets syndrome data and determines the recovery operation or updates the software's record of the state. This repeated measurement loop is one reason fault-tolerant quantum computing requires tight integration between quantum hardware and classical computing.

The surface code is one of the most studied quantum error-correcting codes because it tolerates local interactions on a two-dimensional grid and has a relatively high error threshold. Other codes, including color codes, subsystem codes, bosonic codes, low-density parity-check codes, and hardware-tailored codes, may reduce overhead for particular platforms. The most important engineering milestone is not merely demonstrating one logical qubit, but showing that making the code larger actually suppresses logical error rates and that logical gates can be executed for useful computations.

Error mitigation is different from error correction. Error mitigation uses statistical techniques to estimate what a noisy circuit would have produced with less noise, but it does not protect the quantum state during the computation. Mitigation is useful for current noisy intermediate-scale quantum processors, especially in chemistry and optimization experiments, but it is not enough for Shor-scale cryptanalysis. Fault tolerance requires active error correction that keeps the computation reliable as it runs.

Physical qubits, syndrome checks, and logical qubits

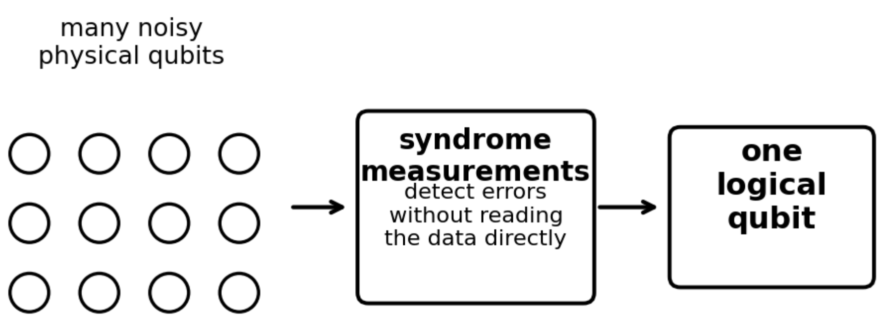

Quantum error correction converts fragile physical states into more reliable logical states.
The price is overhead: many physical qubits and repeated measurements per logical qubit.

Figure 7 Quantum error correction uses many physical qubits and repeated checks to protect one logical qubit.

Table 5 Vocabulary for quantum error correction.

Term	Meaning	Why it matters
Physical qubit	A device-level quantum two-state system.	Transmon, ion, photon, atom, spin, or defect.
Logical qubit	An encoded qubit protected by error correction.	Needed for long, reliable algorithms.
Syndrome measurement	A measurement that detects error information without reading the encoded data.	Repeated continuously during error correction.
Decoder	Classical algorithm that interprets syndrome data.	Turns noisy measurement records into corrections or frame updates.
Threshold	Error rate below which larger codes can reduce logical errors.	A central target for hardware engineering.
Magic state	A prepared resource state used for non-Clifford gates in many fault-tolerant schemes.	Often a major overhead in algorithms.

Benchmarks should be read carefully. Quantum volume, algorithmic qubits, circuit layer operations per second, error per layered gate, heavy-hex or square-lattice connectivity, random circuit sampling results, and application-specific experiments all measure different things. A single headline number rarely tells the whole story. For example, an annealer can have thousands of qubits but cannot run the same gate-model algorithm as a 100-qubit superconducting processor. A photonic sampling machine can use hundreds of modes or thousands of detected photons in a task that is not equivalent to running a universal algorithm. A trapped-ion processor may have fewer physical qubits but higher connectivity and fidelity. Good comparison requires understanding the modality.

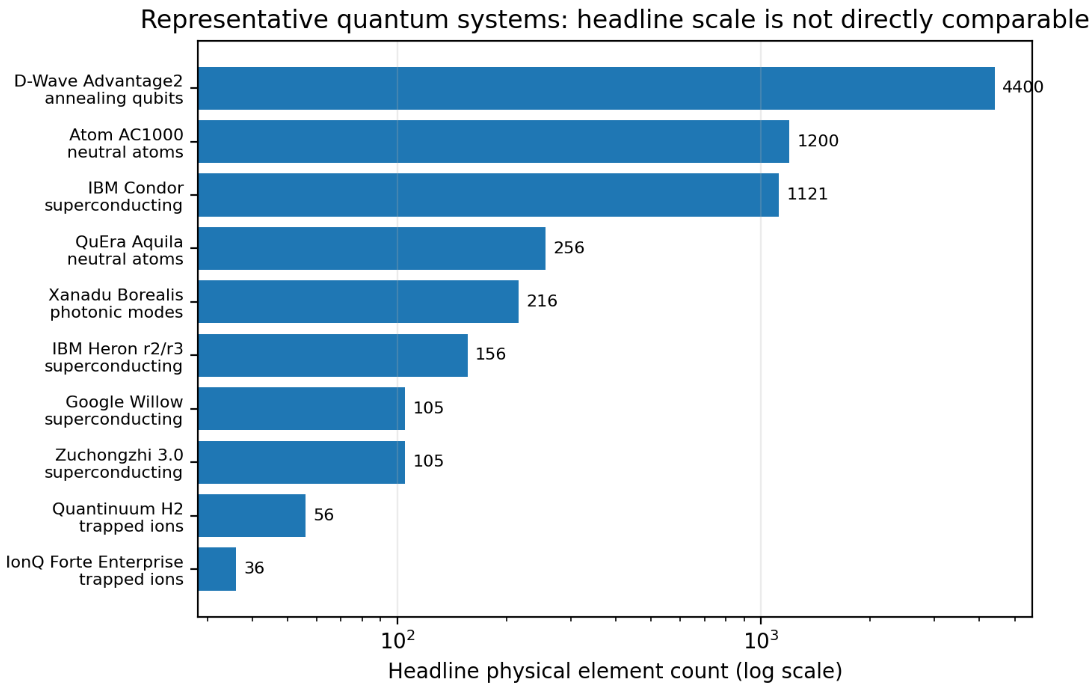

Figure 8 Representative headline element counts. Counts combine qubits, modes, and detected photons and are not directly comparable across hardware types. Sources include IBM, Google, USTC, Quantinuum, IonQ, QuEra, D-Wave, Atom Computing, and Xanadu public materials.

Physical Qubits and Physical Quantum Computers

The abstract qubit is a mathematical object, but a physical quantum computer must embody that object in hardware. Every implementation must solve five core tasks: initialize the qubits, isolate them from unwanted noise, apply controllable gates or evolutions, couple qubits to each other when needed, and measure the final state accurately. Each platform chooses a different compromise among coherence, gate speed, connectivity, fabrication, operating temperature, cost, and scalability.

The term physical quantum computer can refer to several types of machines. A universal gate-model computer is designed to run arbitrary quantum circuits after compilation. An analog quantum simulator is designed to emulate a family of Hamiltonians or quantum dynamics directly. A photonic boson sampler is a specialized sampling device. A quantum annealer is designed for optimization-style energy minimization. These machines all use quantum physics, but they do not have the same computational model or the same cybersecurity implications.

Superconducting quantum computers

Superconducting quantum computers use electrical circuits cooled to extremely low temperatures. The most common qubit type is the transmon, which uses Josephson junctions to create a nonlinear oscillator with two energy levels that can serve as |0> and |1>. Microwave pulses drive gates, resonators assist readout, and lithographic fabrication enables chip-based scaling. IBM, Google, Rigetti, IQM, and several Chinese research groups have used variants of this approach.

The strengths of superconducting circuits are fast gates, compatibility with microfabrication, and the ability to engineer chip layouts. The challenges include coherence, crosstalk, microwave control wiring, cryogenic packaging, fabrication variation, leakage, and calibration drift. Scaling from hundreds to thousands or millions of physical qubits also requires classical control electronics that can operate efficiently near cryogenic environments or communicate reliably through many stages of a dilution refrigerator.

Superconducting processors dominate many public demonstrations because their gates are fast enough to run deep random circuits before the qubits decohere. Google Willow, IBM Heron and Condor, and USTC's Zuchongzhi family are prominent examples. However, the ultimate test is not random circuit sampling alone. The field is moving toward lower error rates, modular packaging, high-quality logical qubits, and sustained circuits useful for chemistry, materials science, optimization, and cryptography.

Trapped-ion quantum computers

Trapped-ion quantum computers store qubits in the internal states of charged atoms held by electromagnetic fields in a vacuum. Lasers or microwaves manipulate the states, and collective motion of the ions can mediate entangling gates. Trapped ions are naturally identical, which avoids some fabrication variability. They often have long coherence times and high-fidelity operations, and small ion chains can provide all-to-all connectivity.

The tradeoff is speed and scaling. Ion gates are typically slower than superconducting gates, and large chains become harder to control. Companies such as Quantinuum use segmented traps and a quantum charge-coupled device architecture to move ions through zones. IonQ has emphasized trapped-ion systems with all-to-all connectivity and high algorithmic fidelity. The central engineering problem is to scale the number of ions while preserving high-quality gates, fast measurement, and practical system operation.

Trapped-ion systems are especially important for error correction research because high fidelity and mid-circuit measurement are valuable. Quantinuum's H2 system, for example, is described as having 56 fully connected qubits, all-to-all connectivity, mid-circuit measurement, qubit reuse, and quantum volume 2^25 (Quantinuum, 2026). IonQ's Forte Enterprise is presented as a 36-qubit data-center-deployable system with single-qubit gate error around 0.02 percent and two-qubit gate error around 0.4 percent (IonQ, 2026).

Photonic quantum computers

Photonic quantum computers use light. A photon qubit can be encoded in polarization, path, time bin, frequency, or other optical degrees of freedom. Photons are attractive because they travel quickly, interact weakly with the environment, and are natural carriers for quantum communication. Photonic systems can operate without the same deep cryogenic requirements as superconducting circuits, although high-efficiency detectors and some sources may still require cooling.

The difficulty is that photons do not naturally interact strongly with each other. Two-qubit gates are therefore hard. Photonic architectures often use measurement-induced nonlinearities, entangled resource states, squeezing, multiplexing, and feed-forward control. Some of the most visible photonic achievements are Gaussian boson sampling experiments, including USTC's Jiuzhang systems and Xanadu's Borealis. These experiments demonstrate difficult sampling tasks but are not the same as general-purpose fault-tolerant quantum computers.

Photonic quantum computing is also central to future quantum networks. If modular quantum processors need to be connected over distance, photons are a natural interface. A scalable quantum internet would rely on photonic links, quantum repeaters, entanglement distribution, and error correction across networks. Thus photonics may be both a computing modality and an interconnect technology for other modalities.

Neutral-atom quantum computers

Neutral-atom quantum computers hold atoms in arrays of optical tweezers, often using rubidium, cesium, strontium, or ytterbium. Qubits can be stored in long-lived internal states. Entangling gates use Rydberg excitation, in which atoms in highly excited states interact strongly over micrometer distances. The Rydberg blockade effect prevents nearby atoms from being simultaneously excited and can implement controlled operations.

The appeal of neutral atoms is scalability and reconfigurability. Optical systems can arrange hundreds or thousands of atoms, move atoms between zones, and perform parallel operations. QuEra's Aquila is described as a 256-qubit analog neutral-atom processor available through Amazon Braket and premium access (QuEra, 2026). Atom Computing's AC1000 product page describes a neutral-atom system with more than 1,200 physical qubits, all-to-all connectivity, mid-circuit measurement, qubit reuse and reset, and real-time conditional branching (Atom Computing, 2026).

Neutral-atom systems must still solve fidelity, addressing, atom loss, laser stability, and error-correction overhead challenges. But they are one of the clearest examples of a platform where raw physical scaling can be rapid. The important question is whether those large arrays can be turned into many high-quality logical qubits that perform useful computations reliably.

Spin qubits and quantum dots

Spin qubits store information in the spin state of an electron, hole, or nucleus, often confined in a semiconductor quantum dot. Silicon, germanium, and gallium arsenide devices have all been studied. The long-term attraction is compatibility with semiconductor manufacturing and the possibility of very dense qubit arrays. A spin qubit can be extremely small compared with an ion trap or a superconducting circuit, which creates an appealing path to scale.

The challenges include fabricating uniform devices, controlling charge noise, coupling distant qubits, and integrating cryogenic electronics. Spin qubits also require careful initialization and readout. Research progress is substantial, but spin-qubit systems are not yet as publicly visible at large scale as superconducting, trapped-ion, neutral-atom, or annealing systems. They remain important because a manufacturing-compatible qubit could become powerful if error rates and interconnects are solved.

Diamond nitrogen-vacancy centers and defect qubits

A nitrogen-vacancy center is a defect in diamond in which a nitrogen atom replaces a carbon atom next to a missing carbon site. The electron spin associated with the defect can be used as a qubit. Such systems can operate under conditions that are attractive for sensing and networking, and they can couple electronic and nuclear spins. Defect qubits in silicon carbide and other materials are also active research areas.

Defect centers are especially promising for quantum sensing and quantum communication nodes because they can combine optical access with spin coherence. They are less mature as large universal quantum computers. The key difficulties are deterministic fabrication, scaling many high-quality defects, integrating photonic structures, and achieving high-fidelity multi-qubit operations. Even if they are not the first platform for large cryptanalytic machines, they may be important for secure networks, sensors, and hybrid quantum devices.

Quantum annealers

Quantum annealers encode optimization problems into an energy function and use quantum dynamics to search for low-energy states. D-Wave is the best-known commercial provider. Its Advantage2 system is described as an annealing quantum computer with more than 4,400 superconducting qubits, a Zephyr topology with 20-way connectivity, and a dilution refrigerator and shielding system (D-Wave, 2026). These systems are not designed to run arbitrary gate-model circuits in the same way as IBM, Google, IonQ, or Quantinuum processors.

That distinction matters for cybersecurity. A quantum annealer with thousands of qubits does not automatically threaten RSA or elliptic-curve cryptography. The cryptographic concern comes from a fault-tolerant gate-model quantum computer capable of running Shor's algorithm at scale. Annealers can still be commercially and scientifically useful for optimization, sampling, machine learning experiments, and hybrid classical-quantum workflows, but they should be evaluated according to the computational model they implement.

Topological qubits

Topological quantum computing attempts to store information in global or nonlocal properties of a system so that local noise has less effect. The most widely discussed path uses Majorana zero modes in engineered topological superconducting systems. Microsoft announced Majorana 1 in 2025 as a topological-core QPU built with topoconductor materials, and described a roadmap toward fault-tolerant prototypes (Microsoft, 2025). The claim is significant because topological protection could reduce error-correction overhead if the physics and engineering scale as intended.

Topological qubits should be treated as a high-potential but still-developing path. Theoretical motivation is strong, but practical validation requires reliable qubit operation, entangling operations, error correction, and independent benchmarking. For an introductory cybersecurity reader, the takeaway is that hardware diversity is increasing. It would be unwise to base a quantum-risk plan on the assumption that only one hardware modality can succeed.

Table 6 Physical platforms for quantum computing.

Platform	Qubit embodiment	Strengths	Key limits
Superconducting transmons	Microwave-controlled energy states in Josephson circuits.	Fast gates, chip fabrication, strong industrial investment.	Cryogenics, coherence, wiring, crosstalk, calibration.
Trapped ions	Internal states of ions in electromagnetic traps.	High fidelity, long coherence, all-to-all connectivity in chains.	Slower gates, optical complexity, scaling ion arrays.
Photonic	Polarization, path, time-bin, frequency, or squeezed optical modes.	Networking compatibility, low decoherence during travel, room-temperature optics possible.	Hard deterministic interactions, source and detector requirements.
Neutral atoms	Internal states of atoms in optical tweezers; Rydberg interactions for gates.	Large arrays, reconfigurability, parallelism, natural uniformity.	Laser control, atom loss, gate fidelity, error-correction overhead.
Spin / quantum dot	Electron, hole, or nuclear spin in semiconductor structures.	Tiny footprint and manufacturing compatibility.	Device variability, cryogenic control, connectivity.
NV and defect centers	Spin states associated with crystal defects.	Useful for sensing and network nodes; possible room-temperature operation.	Scalable fabrication and multi-qubit gates remain difficult.
Quantum annealing	Superconducting flux qubits evolving toward low-energy states.	High qubit counts and optimization-focused workflows.	Not equivalent to universal gate-model computation.
Topological	Nonlocal states such as Majorana modes.	Potentially lower error overhead if demonstrated at scale.	Still experimental and needs broader validation.

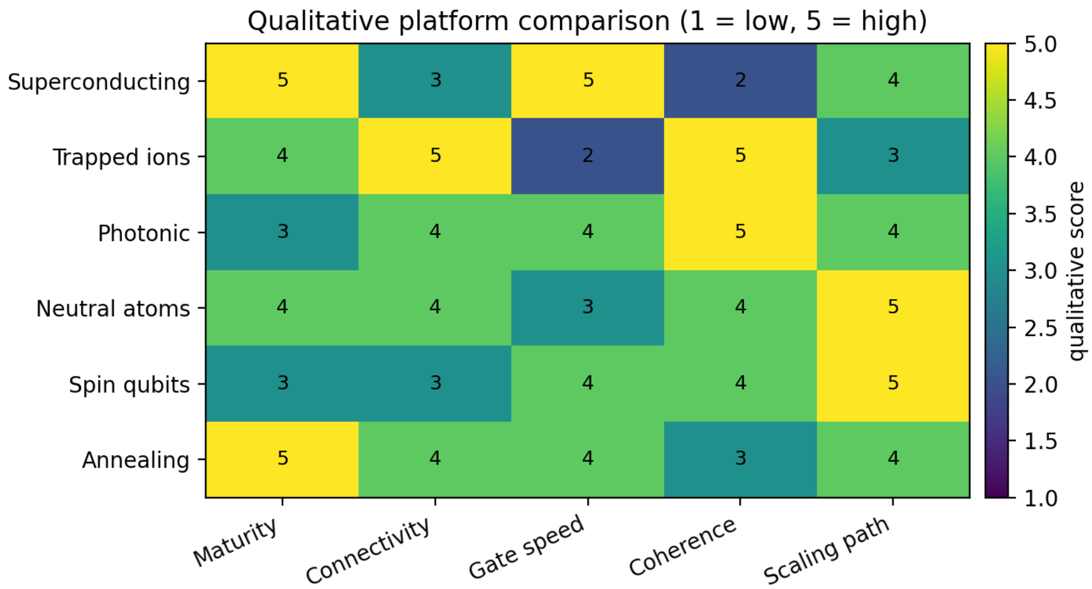

Figure 9 Qualitative comparison of hardware modalities. Scores are conceptual and intended for discussion, not benchmarking.

Major Quantum Computers and Programs

The following survey focuses on prominent systems and programs rather than a comprehensive catalog. The field changes quickly, and announcements often use different metrics. A fair comparison therefore requires noting the modality, the benchmark, and the purpose of the system. A system designed for random circuit sampling, an annealer, a photonic boson sampler, and a trapped-ion gate-model computer are all quantum devices, but their numbers do not mean the same thing.

IBM Quantum

IBM has one of the most visible superconducting quantum-computing programs. Its public roadmap emphasizes a progression from noisy processors toward error-corrected systems. IBM announced the 1,121-qubit Condor processor and the Heron processor family in 2023. IBM describes Heron as using fixed-frequency qubits with tunable couplers to reduce crosstalk, and describes IBM Quantum System Two as a modular system architecture intended for larger quantum processors and multiple chips (IBM, 2023).

IBM's current hardware materials list Eagle, Heron, and Nighthawk families, including Heron r2 and r3 processors with 156 qubits and reported metrics such as error per layered gate and circuit layer operations per second (IBM, 2026a). IBM has also described a roadmap to Starling, targeted for 2029, which the company says is intended to execute 100 million gates over 200 qubits with error correction, and Blue Jay, targeted for 2033, intended to execute one billion gates across 2,000 qubits (IBM, 2023; IBM, 2026b). These roadmap statements are targets, not proof that such machines already exist.

IBM is important for cybersecurity because it explicitly frames the next era around fault tolerance, logical qubits, and useful circuits. High raw qubit counts are only one step. The key cybersecurity milestone would be a machine that can run long algorithms with enough logical qubits and enough reliable gates to threaten deployed public-key cryptography. IBM's roadmap indicates the industry is moving in that direction, but it does not remove the need to migrate now. Cryptographic transition takes many years.

Google Quantum AI

Google's superconducting processors are associated with the Sycamore random circuit sampling experiment and the later Willow chip. In 2019 Google and collaborators reported that Sycamore had performed a sampling task beyond what they estimated was practical for a classical supercomputer. The claim generated debate because classical simulation methods continued to improve, but it remains an important milestone in the experimental history of quantum computing.

Google announced Willow in 2024 and later updated its public description in 2025. Google describes Willow as a 105-qubit superconducting chip that achieved two notable results: exponential error suppression as the surface-code array was scaled from 3 by 3 to 5 by 5 to 7 by 7, and a random circuit sampling benchmark completed in under five minutes that Google estimated would take the fastest supercomputer 10 septillion years under stated assumptions (Google Quantum AI, 2025). Google also emphasizes that random circuit sampling has no known commercial application and that the next challenge is useful beyond-classical computation with real-world relevance.

For cybersecurity readers, Google Willow is important less because it breaks cryptography and more because it points toward the threshold behavior required for error correction. Showing that logical errors can be reduced by increasing code size is one of the bridges from noisy processors to useful fault-tolerant machines. The gap from that milestone to attacking RSA or ECC is still large, but the direction is clear.

China: Zuchongzhi and Jiuzhang

Chinese research groups, especially teams associated with the University of Science and Technology of China, have reported major superconducting and photonic quantum advantage demonstrations. Zuchongzhi is a family of superconducting processors used for random circuit sampling. A 2024 preprint on Zuchongzhi 3.0 describes a 105-qubit superconducting processor with reported single-qubit, two-qubit, and readout fidelities of 99.90 percent, 99.62 percent, and 99.18 percent, and reports an 83-qubit, 32-cycle random circuit sampling experiment with a claimed large classical cost estimate (Gao et al., 2024).

Jiuzhang is a photonic system used for Gaussian boson sampling. Jiuzhang 2.0 was described as a phase-programmable photonic quantum computer with up to 113 detected photons out of a 144-mode circuit and a sampling rate claimed to be far beyond brute-force classical simulation (Zhong et al., 2021). Later reports on Jiuzhang 4.0 describe very large programmable photonic boson-sampling experiments, although, as with all fast-moving preprints, benchmark claims should be read with attention to assumptions and peer review.

The Chinese demonstrations matter because they show that quantum advantage claims are not limited to a single company or a single physical platform. They also illustrate the difference between special-purpose sampling experiments and general-purpose fault tolerance. Zuchongzhi and Jiuzhang are scientifically significant, but cybersecurity impact depends on future machines that can execute long, error-corrected algorithms.

IonQ and Quantinuum

IonQ and Quantinuum are leading trapped-ion companies. IonQ emphasizes commercial access, data-center deployment, and trapped-ion connectivity. Its Forte Enterprise page describes a 36-qubit system, high single-qubit and two-qubit fidelities, and coherence times that are much longer than typical superconducting coherence times (IonQ, 2026a). IonQ also describes Tempo as a future-oriented system with a target of 100 qubits, faster gate speeds, mid-circuit measurement, and 99.9 percent target fidelity (IonQ, 2026b).

Quantinuum's H-Series processors use trapped ions and a quantum charge-coupled-device architecture. The H2 page describes a 56-qubit system with all-to-all connectivity, mid-circuit measurement, conditional logic, qubit reuse, quantum volume 2^25, greater than 99.99 percent single-qubit fidelity, and greater than 99.9 percent two-qubit fidelity (Quantinuum, 2026). Research reports in 2025 also describe Helios, a 98-qubit trapped-ion quantum computer using a QCCD architecture (Pino et al., 2025).

The main importance of trapped-ion systems is quality. Fewer physical qubits with high connectivity and low error can be more useful than more qubits with high noise for some experiments. Trapped-ion machines are also strong candidates for early logical-qubit and error-correction demonstrations because mid-circuit measurement and qubit reuse are native strengths.

The open question is how quickly ion-trap architectures can scale while maintaining those advantages.

Photonic companies: Xanadu and others

Xanadu's Borealis experiment is one of the best-known photonic demonstrations. The 2022 Nature paper describes a programmable photonic processor that performed Gaussian boson sampling on 216 squeezed modes. The paper reported that Borealis produced samples in 36 microseconds and estimated that the best classical exact method would require more than 9,000 years per sample under the assumptions in the paper (Madsen et al., 2022).

Photonic companies and research groups pursue several paths. Some focus on boson sampling and special-purpose quantum advantage. Others aim at universal measurement-based quantum computation using large photonic cluster states. Still others emphasize quantum networking and interconnects. Photonic hardware is important because photons are the natural communication medium for distributed quantum systems, even if another modality becomes dominant for local processing.

Neutral-atom companies: QuEra, Atom Computing, Pasqal, and others

Neutral-atom companies are attracting attention because optical tweezer arrays can scale to large numbers of atoms. QuEra's Aquila is presented as a 256-qubit analog neutral-atom quantum computer, available through Amazon Braket and premium access. Its public materials describe dynamic qubit layout, analog programming, and applications in simulation, optimization, and machine learning workflows (QuEra, 2026).

Atom Computing's AC1000 page describes more than 1,200 physical qubits based on ytterbium-171 nuclear-spin qubits, all-to-all connectivity, mid-circuit measurement with qubit reset and reuse, real-time conditional branching, greater than 99.9 percent one-qubit gate fidelity, greater than 99.6 percent two-qubit gate fidelity, and greater than 99.8 percent SPAM fidelity (Atom Computing, 2026). These are ambitious product-level claims and should be interpreted in the context of the system's actual availability, benchmarking, and use cases.

Pasqal and other neutral-atom groups also pursue analog and digital neutral-atom quantum processors. The shared engineering themes are laser control, atom rearrangement, Rydberg gates, and modular scaling. Neutral atoms may prove important for early logical-qubit demonstrations because they can provide many physical qubits in reconfigurable geometries.

D-Wave and quantum annealing

D-Wave's systems are often cited because their qubit counts are large. The company's Advantage2 system is described as containing more than 4,400 superconducting qubits, using Zephyr topology with 20-way connectivity, and supporting on-premises or cloud use through Leap (D-Wave, 2026). These systems are built for annealing and hybrid optimization workflows, not for general gate-model algorithms. The qubit count is therefore not comparable to IBM, Google, IonQ, or Quantinuum qubit counts.

Annealing machines can be valuable for optimization, logistics, scheduling, sampling, and hybrid workflows. They are also useful in the broader commercial ecosystem because they provide organizations with real quantum hardware experience. However, they are not the direct reason RSA, Diffie-Hellman, and elliptic-curve cryptography must migrate. The direct reason is the prospect of fault-tolerant universal quantum computers.

Table 7 Major quantum computing programs and representative systems.

Organization	Representative system	Modality	Important context
IBM	Condor; Heron; Nighthawk; System Two; Starling roadmap	Superconducting	Condor 1,121 qubits; Heron family up to 156 qubits; roadmap to Starling and Blue Jay fault-tolerant systems.
Google	Sycamore; Willow	Superconducting	Willow described as 105 qubits with below-threshold surface-code scaling and a major RCS benchmark.
USTC / China	Zuchongzhi; Jiuzhang	Superconducting and photonic	Zuchongzhi 3.0 reports 105 superconducting qubits; Jiuzhang systems target Gaussian boson sampling.
Quantinuum	H2; Helios reported	Trapped ion	H2 described as 56 fully connected qubits with high fidelities and QV 2^25.
IonQ	Forte Enterprise; Tempo target	Trapped ion	Forte Enterprise described as 36 qubits; Tempo target 100 qubits.
Xanadu	Borealis	Photonic	216 squeezed modes used for Gaussian boson sampling demonstration.
QuEra	Aquila	Neutral atom	256-qubit analog neutral-atom processor with cloud access.
Atom Computing	AC1000	Neutral atom	More than 1,200 physical qubits and logical-qubit-era features claimed on product page.
D-Wave	Advantage2	Quantum annealing	More than 4,400 superconducting annealing qubits; not a gate-model machine.
Microsoft	Majorana 1 roadmap	Topological research path	Announced topological-core QPU roadmap; still requires validation and scaling.

Source note: Values in this table are drawn from public company pages and papers cited in the References. Modalities use different metrics; do not compare rows by count alone.

Interpreting quantum computer claims

Quantum computing is full of impressive numbers, but cybersecurity professionals must learn to interpret them. Qubit count matters only when combined with fidelity, connectivity, gate speed, measurement quality, calibration stability, and error-correction overhead. A processor with many noisy qubits may be less useful than a smaller processor with excellent gates. A benchmark that is strong for one modality may not apply to another.

Random circuit sampling is a benchmark in which a quantum processor produces samples from the output distribution of a random quantum circuit. It is useful because the distribution becomes hard to simulate classically as circuit size and depth increase. However, RCS is not itself a business application or a cryptographic attack. Gaussian boson sampling is similar in that it can be classically hard under certain assumptions, but it is a specialized photonic sampling task. Quantum volume and algorithmic qubits aim to combine size and quality, but they also have limitations.

For a cybersecurity roadmap, the key questions are different. How many logical qubits can the system maintain? What is the logical error rate? How many reliable logical gates can it perform? Does it support the non-Clifford operations needed for universal fault-tolerant computing? How long does error correction remain stable? How much classical processing is required for decoding? Does the machine have enough logical width and depth to run resource estimates for breaking RSA-2048 or elliptic-curve keys? These are the questions that connect hardware progress to cryptographic risk.

A second issue is time horizon. Hardware roadmaps are not guarantees. They are engineering plans. Some targets will slip; some approaches will fail; some unexpected approach may accelerate. For that reason, the rational security stance is not panic, but preparation. Inventory cryptographic dependencies, protect long-lived data, test post-quantum protocols, and build agility before a crisis forces rapid migration.

Table 8 How to read quantum-computer benchmarks.

Metric	What it means	Caution
Qubit count	Number of physical qubits, atoms, ions, or other elements.	Does not measure quality or logical capability.
Gate fidelity	How often a gate behaves as intended.	Reported methods vary; two-qubit fidelity is often the bottleneck.
Coherence time	How long a qubit preserves useful state information.	Must be compared with gate speed and circuit depth.
Connectivity	Which qubits can interact directly.	Affects circuit overhead and compilation cost.
Quantum volume	Benchmark combining size and error for certain circuits.	Useful but not universal; not a direct application metric.
Random circuit sampling	Sampling benchmark for hard-to-simulate random circuits.	Demonstrates special-purpose advantage, not general utility by itself.
Logical qubit error rate	Error rate after quantum error correction.	Most relevant for fault-tolerant applications.
CLOPS / throughput	Circuit layer operations per second.	Important for user throughput and iterative workloads.

Applications beyond cryptography

Although cybersecurity receives attention because of Shor's algorithm, many quantum-computing applications are positive. The most natural application is simulation of quantum systems. Chemistry, catalysis, materials science, battery design, superconductors, nitrogen fixation, and drug discovery all involve quantum behavior that is difficult to model exactly on classical computers. A fault-tolerant quantum computer could estimate molecular energies and reaction pathways with high precision, potentially reducing expensive laboratory search.

Optimization is another major area, but it is often misunderstood. Many optimization problems are NP-hard, and quantum computers are not expected to solve all NP-hard problems efficiently. Instead, the hope is that quantum heuristics, annealing, or hybrid algorithms may provide speedups or better solutions for specific structured instances. Logistics, scheduling, portfolio optimization, power-grid planning, and manufacturing are common examples. Evidence for broad practical advantage remains limited, but the application area motivates experimentation.

Machine learning is also frequently discussed. Quantum machine learning may provide advantages in data representation, kernel estimation, sampling, or training certain models, but claims must be evaluated carefully because loading classical data into a quantum state can be expensive. Near-term quantum machine learning is largely exploratory. Longer-term, quantum computers may be useful when the data or model is natively quantum, such as in quantum chemistry or quantum sensor networks.

Quantum sensing and quantum communication are related but distinct fields. Quantum sensors use superposition and entanglement to measure fields, time, acceleration, or material properties with high precision. Quantum communication uses quantum states to distribute keys or entanglement. Quantum key distribution can detect eavesdropping under specific assumptions, but it does not replace the need for post-quantum cryptography in ordinary internet protocols. A future quantum network may combine photonic links, memories, repeaters, and small processors to support distributed quantum computing.

Table 9 Major application areas beyond cryptanalysis.

Application area	Examples	Status
Quantum simulation	Chemistry, materials, condensed matter	Probably one of the most natural long-term applications.
Optimization	Routing, scheduling, resource allocation	Likely hybrid and problem-specific; no universal magic optimizer.
Sampling	Random circuit sampling, boson sampling, generative tasks	Useful for benchmarking and some research; application value varies.
Machine learning	Kernels, sampling, quantum data, hybrid models	Promising but data-loading and advantage questions remain.
Sensing	Magnetometry, clocks, inertial sensing, materials characterization	Often nearer-term than large fault-tolerant computing.
Networking	Entanglement distribution and quantum repeaters	Needed for modular and distributed quantum systems.

Possible quantum resistant cryptographic algorithms

The primary reason quantum computing is a concern for cybersecurity professionals is that quantum algorithms already exist for the mathematical problems underlying current public-key cryptography. Post-quantum cryptography seeks algorithms that run on ordinary classical computers but resist both classical and quantum attacks. The goal is not to use quantum hardware for encryption. The goal is to deploy classical algorithms whose security does not rest on factoring or discrete logarithms.

Several mathematical families have been studied. Lattice-based cryptography is the leading family for general-purpose key establishment and many digital signatures. Code-based cryptography, hash-based signatures, multivariate cryptography, and isogeny-based cryptography have also been investigated. The NIST process selected algorithms after years of public analysis, but cryptography never has absolute certainty. Algorithms can be broken, parameters can be adjusted, and implementation flaws can undermine strong mathematics. This is why cryptographic agility remains important.

Lattice-based cryptography uses problems involving high-dimensional lattices. A lattice is a regular set of points generated by integer combinations of basis vectors. Problems such as the Shortest Vector Problem, Closest Vector Problem, Learning With Errors, Module Learning With Errors, and Ring Learning With Errors are believed to be hard for both classical and quantum computers at appropriate parameters. ML-KEM and ML-DSA are lattice-based standards derived from the CRYSTALS family. Their efficiency and relatively small key sizes compared with some alternatives make them practical for many protocols.

Hash-based signatures rely on the security of cryptographic hash functions rather than number-theoretic assumptions. SLH-DSA, derived from SPHINCS+, is stateless, meaning it avoids the state-management hazards of older one-time or few-time hash-signature schemes. Its signatures are larger and often slower than lattice signatures, but its conservative assumptions make it valuable as a backup standard. For some high-assurance environments, algorithm diversity may be as important as raw performance.

Code-based cryptography dates back to the McEliece cryptosystem, proposed in 1978. It has withstood decades of study, but public keys are large. Multivariate cryptography builds on systems of polynomial equations over finite fields, but several proposals have been broken. Isogeny-based cryptography attracted attention because of small keys, but the SIKE candidate was broken during the NIST process. These outcomes are not failures of the process; they are evidence that open cryptanalysis is necessary before deployment.

Migration should be staged. First, inventory where cryptography is used. Second, classify assets by how long confidentiality and authenticity must last. Third, prioritize internet-facing protocols, certificate systems, code signing, firmware updates, VPNs, and long-lived sensitive data. Fourth, test hybrid deployments and post-quantum libraries in realistic environments. Fifth, update policies, procurement language, and incident-response plans. Finally, maintain the ability to change algorithms again if cryptanalysis or standards evolve.

Table 10 NIST post-quantum cryptography standards and planned signature work.

Standard	Algorithm name	Original family	Use	Notes
FIPS 203	ML-KEM	CRYSTALS-Kyber	Key establishment / general encryption support	Primary standardized KEM for many protocols.
FIPS 204	ML-DSA	CRYSTALS-Dilithium	Digital signatures	Primary lattice-based signature standard.
FIPS 205	SLH-DSA	SPHINCS+	Digital signatures	Stateless hash-based signature backup.
Future / draft path	FN-DSA	Falcon	Digital signatures	NIST has indicated a Falcon-derived standard is planned.

Source note: NIST released FIPS 203, FIPS 204, and FIPS 205 in 2024 and advised organizations to begin integrating them rather than waiting for a future quantum attack.

Table 11 Practical quantum-readiness migration checklist.

Step	Action
1. Inventory	Find TLS, VPN, SSH, PKI, code signing, firmware update, backups, databases, and embedded devices that use public-key cryptography.
2. Classify data lifetime	Identify information that must remain confidential for years or decades and is exposed to harvest-now-decrypt-later risk.
3. Test hybrid protocols	Combine classical and post-quantum key exchange during transition to reduce dependence on any single assumption.
4. Update certificates	Plan for larger public keys and signatures in certificate chains, hardware security modules, and directory services.
5. Measure performance	Benchmark handshake size, latency, memory use, CPU load, and failure behavior under production conditions.
6. Build agility	Use libraries, protocols, and policies that can change algorithms and parameters without system redesign.
7. Monitor standards	Track NIST, IETF, vendor guidance, and cryptanalytic developments.

Security planning for the quantum transition

A useful planning assumption is that migration will take longer than expected. Many organizations still depend on old TLS configurations, legacy VPNs, embedded devices with hard-coded trust anchors, certificate authorities with long operational cycles, and hardware that cannot be updated easily. A quantum-safe migration can be slowed by interoperability, packet sizes, hardware acceleration, regulatory approvals, procurement timelines, and user devices outside organizational control.

The most urgent category is long-lived confidentiality. If an adversary can record encrypted traffic today and decrypt it later, then today's data is at risk even before a quantum computer exists. Examples include diplomatic cables, intelligence records, health records, biometric data, intellectual property, design files, and long-term financial records. For these assets, the relevant deadline is not the day a quantum computer breaks RSA. It is the last day on which vulnerable cryptography protects data that must remain secret into the quantum era.

Authenticity also matters. Digital signatures protect software updates, certificates, legal documents, audit logs, and identity assertions. An attacker with future quantum capabilities could forge signatures under old keys if those keys remain trusted. Long-lived signatures may need timestamping, archival validation, and transition plans. Secure boot and firmware update ecosystems are especially important because weak signature migration can create durable supply-chain vulnerabilities.

Finally, organizations should avoid replacing one monoculture with another. It is prudent to prefer standardized algorithms, but it is also prudent to maintain agility. Post-quantum cryptography is still young compared with RSA and ECC deployment history. Strong engineering will include hybrid modes, defense in depth, implementation hardening, side-channel resistance, formal verification where possible, and monitoring for new cryptanalytic results.

Conclusions

Quantum computers store and process information in quantum states. A qubit is not merely a faster bit; it is a two-level quantum system described by amplitudes, phase, and measurement probabilities. Multi-qubit systems can be entangled, and quantum gates manipulate amplitudes so that algorithms can exploit interference. This model gives large speedups for some structured problems, including the problems that protect many public-key cryptosystems.

Physical quantum computers are diverse. Superconducting circuits, trapped ions, photons, neutral atoms, semiconductor spins, defect centers, annealers, and topological approaches each make different engineering tradeoffs. Superconducting systems are fast and chip-based but cryogenic and noisy. Trapped ions are high fidelity and highly connected but harder to scale quickly. Photons are excellent carriers of quantum information but difficult to make interact. Neutral atoms scale rapidly in optical arrays but must convert scale into reliable logic. Annealers target optimization rather than universal gate-model computation. Topological qubits offer an intriguing path to lower error overhead but remain a developing technology.

Major programs from IBM, Google, Chinese research groups, IonQ, Quantinuum, Xanadu, QuEra, Atom Computing, D-Wave, Microsoft, and others show that the field is moving quickly. Yet the cybersecurity impact depends on fault-tolerant logical qubits and long reliable computations, not raw physical qubit counts alone. Today's demonstrations are important scientific and engineering milestones, but they do not mean that current public-key cryptography has already failed in practice.

The correct cybersecurity response is preparation. NIST has finalized initial post-quantum standards, and organizations should begin inventory, testing, hybrid deployment, certificate planning, and cryptographic agility work. The cost of migrating calmly over several years is much lower than the cost of replacing foundational security protocols under emergency conditions. Quantum computing is both an opportunity and a threat; cybersecurity professionals must understand the physics well enough to plan realistically, without either panic or complacency.

References

Atom Computing. (2026). AC1000 product information. Atom Computing public product page.

D-Wave. (2026). The Advantage2 quantum computer. D-Wave public systems page.

Gao, X., et al. (2024). Establishing a new benchmark in quantum computational advantage with 105-qubit Zuchongzhi 3.0 processor. arXiv:2412.11924.

Google Quantum AI. (2025). Meet Willow, our state-of-the-art quantum chip. Google Research Blog, updated June 12, 2025.

Grover, L. K. (1996). A fast quantum mechanical algorithm for database search. Proceedings of the ACM Symposium on Theory of Computing.

IBM. (2023). IBM Quantum roadmap to 2033, including Condor, Heron, System Two, Starling, and Blue Jay announcements. IBM Quantum Blog.

IBM. (2026a). IBM Quantum hardware public information page. IBM Quantum.

IBM. (2026b). IBM commits more than $10 billion over five years to build a fault-tolerant quantum computer. IBM Quantum Blog.

IonQ. (2026a). IonQ Forte Enterprise technical information. IonQ public system page.

IonQ. (2026b). IonQ Tempo technical information. IonQ public system page.

Lavor, C., Manssur, L. R. U., & Portugal, R. (2003). Shor's algorithm for factoring large integers. arXiv quantum computing tutorial literature.

Madsen, L. S., et al. (2022). Quantum computational advantage with a programmable photonic processor. Nature, 606, 75-81.

Microsoft. (2025). Microsoft unveils Majorana 1, the world's first quantum processor powered by topological qubits. Microsoft Azure Quantum Blog.

Nielsen, M. A., & Chuang, I. L. (2002). Quantum Computation and Quantum Information. Cambridge University Press.

NIST. (2024). NIST releases first three finalized post-quantum encryption standards: FIPS 203, FIPS 204, and FIPS 205. National Institute of Standards and Technology.

Pino, J. M., et al. (2025). Helios: A 98-qubit trapped-ion quantum computer. arXiv:2511.05465.

Preskill, J. (2018). Quantum Computing in the NISQ era and beyond. Quantum, 2, 79.

Quantinuum. (2026). System Model H2 technical information. Quantinuum public hardware page.

QuEra. (2026). Aquila: 256-qubit analog neutral-atom quantum computer. QuEra public product page.

Rieffel, E., & Polak, W. (2011). Quantum Computing: A Gentle Introduction. MIT Press.

Shor, P. W. (1994). Algorithms for quantum computation: Discrete logarithms and factoring. Proceedings of the IEEE Symposium on Foundations of Computer Science.

Wurtz, J., et al. (2023). Aquila: QuEra's 256-qubit neutral-atom quantum computer. arXiv:2306.11727.

Zhong, H.-S., et al. (2021). Phase-programmable Gaussian boson sampling using stimulated squeezed light. arXiv:2106.15534.